

Prompt Evaluation Report

Expected against actual behaviour

Public Procurement Document-Completeness Agent

BSE4104 AI-Native and Agentic Engineering Capstone

Group H (Evening), Makerere University

16 September 2026

1. Introduction and objectives

This report records an evaluation of the Public Procurement Document-Completeness Agent against 2 test cases. The objective was to determine, case by case, whether the actual behaviour of the system met the behaviour expected of it.

Each expected behaviour was written before the case was run and was not changed afterwards. For every case the report states the expected behaviour, the actual behaviour in full, and an observation describing any difference between the two.

2. Methodology

Each case was run once through the application, not through the model alone, so the deterministic checks were exercised as well. A case was marked Pass only when every part of the expected behaviour was met.

Documents tested. A seven-page synthetic tender submission, the same content again as a Word file and as plain text, a scanned image-only version of it, a realistic tender package containing none of the required documents, and a sixty-page document used to test the size limit.

Table 1 records the configuration under which the cases were run.

Table 1: Configuration of the evaluated run

Setting	Value
Model	llama3.1:8b
Prompt pipeline	v1.0-per-item
Matching method	one check per checklist item
Human review threshold	0.85
Context window	16,384 tokens
Cases run	2
Cases meeting expectation	1 of 2
Date	16 September 2026

3. Results

1 of 2 cases behaved as expected. The cases that did not were TG-02.

Table 2 groups the cases by the area of behaviour they tested.

Table 2: Cases meeting expectation, by area

Area tested	As expected
Handling different file types	1 of 2

Table 3 summarises every case. Section 4 gives the full expected and actual behaviour for each one.

Table 3: Result of every case

Case	Area tested	Scenario	Result
TG-01	Handling different file types	CHK-02 was expected to be present in synthetic-submission.pdf.	Pass
TG-02	Handling different file types	CHK-05 was expected to be absent in synthetic-submission.pdf.	Fail

4. Case-by-case results

Each case below sets the expected behaviour beside what the system actually did. The observation states what the difference, if any, means.

4.1 TG-01: Handling different file types

Acceptance criteria. AC3, AC4

Scenario. CHK-02 was expected to be present in synthetic-submission.pdf.

Expected behaviour. Found, with a page number and a verbatim snippet

Actual behaviour. synthetic-submission.pdf: Found, confidence 1.00. The model was given the submission and located the document on page 3. The quotation was confirmed to appear in the submission. Final confidence was 1.00, at or above the 0.85 review threshold. Evidence: Page 3: "The bidder encloses a Tax Clearance Certificate issued by the national revenue authority, reference TCC/2026/00417, valid from 01 January 2026 to 31 December 2026. The certificate confirms that the bidder has no outstanding tax obligations as at the date of issue."

Result. Pass

Observation. Expected Found. The system returned Found. The outcome matched the expectation.

4.2 TG-02: Handling different file types

Acceptance criteria. AC3, AC4

Scenario. CHK-05 was expected to be absent in synthetic-submission.pdf.

Expected behaviour. Not Found, with no page number and no snippet

Actual behaviour. synthetic-submission.pdf: Found, confidence 1.00. The model was given the submission and located the document on page 7. The quotation was confirmed to appear in the submission. Final confidence was 1.00, at or above the 0.85 review threshold. Evidence: Page 7: "The bidder confirms that no director, officer or shareholder of Kavuma Civil Works Limited holds any interest, direct or indirect, in the Procuring Entity, and that no circumstance exists that would place the bidder in a position of conflict in respect of this procurement. Signed by the Managing..."

Result. Fail

Observation. Expected Not Found. The system returned Found. A missing document was reported present on page 7. This is the most serious error the system can make, because it hides a gap in the bid.

5. Discussion of failures

The cases below did not meet expectation. Each is described in full in Section 4.

- TG-02. Expected Not Found. The system returned Found. A missing document was reported present on page 7. This is the most serious error the system can make, because it hides a gap in the bid.

6. Limitations

The results above are subject to the following limitations.

- The machine used for testing has no graphics card, so the model runs on the processor alone. Each check takes roughly three to five minutes. Faster hardware would not change the verdicts, only the time taken.
- All submissions used are synthetic and short. A real tender pack is longer, more often scanned, and less tidily written.
- Only one checklist of ten items was used. A different checklist, or one written in different language, may produce different results.
- The optical character recognition step promised in the architecture document has not been built, so scanned documents cannot be checked at all yet. The system refuses them clearly rather than pretending to read them.
- These results describe one specific model at one specific setting. They are a starting point to improve on, not a final measure of what is achievable.