

Evaluation of no-items-submission-pdf

Expected against actual behaviour

Public Procurement Document-Completeness Agent

BSE4104 AI-Native and Agentic Engineering Capstone

Group H (Evening), Makerere University

16 September 2026

1. Introduction and objectives

This report records an evaluation of the Public Procurement Document-Completeness Agent against 1 test case. The objective was to determine, case by case, whether the actual behaviour of the system met the behaviour expected of it.

Each expected behaviour was written before the case was run and was not changed afterwards. For every case the report states the expected behaviour, the actual behaviour in full, and an observation describing any difference between the two.

2. Methodology

Each case was run once through the application, not through the model alone, so the deterministic checks were exercised as well. A case was marked Pass only when every part of the expected behaviour was met.

Documents tested. The submission no-items-submission-pdf, one of the documents used for testing.

Table 1 records the configuration under which the cases were run.

Table 1: Configuration of the evaluated run

Setting	Value
Model	llama3.1:8b
Prompt pipeline	v2.0-per-item + adjudicator-v1.0
Matching method	one check per checklist item
Human review threshold	0.85
Context window	16,384 tokens
Cases run	1
Cases meeting expectation	1 of 1
Date	16 September 2026

3. Results

1 of 1 cases behaved as expected.

Table 2 groups the cases by the area of behaviour they tested.

Table 2: Cases meeting expectation, by area

Area tested	As expected
Noticing documents that are missing	1 of 1

Table 3 summarises every case. Section 4 gives the full expected and actual behaviour for each one.

Table 3: Result of every case

Case	Area tested	Scenario	Result
EV-18	Noticing documents that are missing	A complete, realistic tender package that contains NONE of the ten required documents: only a method statement, programme, plant schedule, personnel list and safety approach. Three unrelated checklist items are checked against it. This isolates whether the model can report absence at all, or only ever latches onto the nearest plausible text.	Pass

4. Case-by-case results

Each case below sets the expected behaviour beside what the system actually did. The observation states what the difference, if any, means.

4.1 EV-18: Noticing documents that are missing

Acceptance criteria. AC4

Scenario. A complete, realistic tender package that contains NONE of the ten required documents: only a method statement, programme, plant schedule, personnel list and safety approach. Three unrelated checklist items are checked against it. This isolates whether the model can report absence at all, or only ever latches onto the nearest plausible text.

Document tested. no-items-submission-pdf

Expected behaviour. All three items Not Found, with no page number and no snippet

Actual behaviour. 3 of 3 items came back as expected.

- **CHK-01** (Certificate of Incorporation or business registration certificate): Not Found, confidence 0.00. The model was given all 5 pages of the submission and asked to find this document. It reported that no passage is this document, and returned no page and no quotation. Its match confidence was 0.00. A value of 0.00 means it saw no likely match. With nothing quoted, there was no evidence for the later checks to verify. Evidence: None. No page number and no quotation were returned.
- **CHK-02** (Valid tax clearance certificate from the revenue authority): Not Found, confidence 0.00. The model was given all 5 pages of the submission and asked to find this document. It reported that no passage is this document, and returned no page and no quotation. Its match confidence was 0.00. A value of 0.00 means it saw no likely match. With nothing quoted, there was no evidence for the later checks to verify. Evidence: None. No page number and no quotation were returned.
- **CHK-04** (Bid security in the form and amount stated in the bidding document): Not Found, confidence 0.00. The model was given all 5 pages of the submission and asked to find this document. It reported that no passage is this document, and returned no page and no quotation. Its match confidence was 0.00. A value of 0.00 means it saw no likely match. With nothing quoted, there was no evidence for the later checks to verify. Evidence: None. No page number and no quotation were returned.

Result. Pass

Observation. CHK-01: Expected Not Found. The system returned Not Found. The outcome matched the expectation. CHK-02: Expected Not Found. The system returned Not Found. The outcome matched the expectation. CHK-04: Expected Not Found. The system returned Not Found. The outcome matched the expectation.

5. Limitations

The results above are subject to the following limitations.

- The machine used for testing has no graphics card, so the model runs on the processor alone. Each check takes roughly three to five minutes. Faster hardware would not change the verdicts, only the time taken.
- All submissions used are synthetic and short. A real tender pack is longer, more often scanned, and less tidily written.
- Only one checklist of ten items was used. A different checklist, or one written in different language, may produce different results.
- The optical character recognition step promised in the architecture document has not been built, so scanned documents cannot be checked at all yet. The system refuses them clearly rather than pretending to read them.
- These results describe one specific model at one specific setting. They are a starting point to improve on, not a final measure of what is achievable.