

Completeness Agent: Test Results

Public Procurement Document-Completeness Agent

BSE4104 Capstone, Group-H (Evening), Makerere University

Generated 14 September 2026, 01:03 UTC

Model tested: llama3.1:8b | Prompt version: v1.0-per-item | Method: one check per checklist item

Result: Every check behaved as expected.

What this document is

This is a record of testing carried out on a software assistant that checks whether a tender submission contains every document a procurement checklist requires. It is written to be readable without any technical background.

What the assistant is for

When a supplier submits a bid, that bid must contain a set of required documents: a certificate of incorporation, a tax clearance certificate, a bid security, and so on. Today a human reviewer reads the whole package and ticks each item off a list by hand. That is slow, and an item can be missed when it appears under an unexpected heading or deep in an attachment.

The assistant reads the checklist and the submission and reports, for each required item, whether it could find it, on which page, and quoting the exact words it relied on. A human then verifies the findings.

What the assistant must never do

The assistant is deliberately limited. It must never score or grade a bid, never rank or compare bidders, never comment on whether a clause is legally valid, and never recommend awarding or rejecting a contract. Those judgements belong to authorised procurement officers. Several of the tests below exist purely to confirm the assistant refuses to cross that line.

How to read the results

Each test is one scenario. Before running anything, we wrote down what *should* happen. We then ran it and recorded what *actually* happened. A test passes only if the two match.

A failed test is useful information, not a disaster. This is a first working version, and the purpose of testing it is to find out exactly where it falls short so the next version can fix those specific things.

What was tested against

The submission scanned-submission, one of the documents used for testing.

Every document used is invented. No real bidder, no real procuring entity, and no confidential information was used at any point.

The findings in plain English

Handling different file types and bad files: 1 of 1 behaved as expected.

The assistant was given Word files, plain text, a scanned document with no readable text, and a document too large to process.

Every test, one by one

Handling different file types and bad files

EV-17. A scanned, image-only PDF with no text layer, which the Project Charter names as a document-quality risk. The Architecture and Context Diagram promises an OCR fallback; it does not exist yet, so this case records the real behaviour.

We expected: Rejected with a clear error telling the user OCR is needed. It must fail loudly, never silently report every item as missing.

What happened: Rejected with a clear error: scanned-submission.pdf contains no extractable text. If it is a scanned document, it must be passed through OCR before it can be checked.

Behaved as expected · took instant

Terms used in this report

- **Checklist item.** One document the submission is required to contain, for example a tax clearance certificate.
- **Found / Not Found / Requires Human Review.** The three verdicts the assistant can give. The third means it is unsure and a person must look.
- **Confidence.** A number from 0 to 1 the assistant reports alongside each verdict, meant to express how sure it is. Below 0.85, the item is sent for human review automatically.
- **Evidence snippet.** The exact words copied from the submission that the assistant says prove an item is present. The system checks that this text really appears in the document, so a quotation can never be invented.
- **Model.** The artificial intelligence component, here llama3.1:8b, running privately on a local machine rather than on the internet, so no tender data leaves the computer.
- **Context window.** How much text the model can read at once, here 16,384 units of text. Anything larger is refused outright rather than being quietly cut short, because a cut-short document would make present items look missing.
- **Prompt.** The written instructions given to the model telling it what job to do and what it must never do.

Limitations of this testing

- The machine used for testing has no graphics card, so the model runs on the processor alone. Each check takes roughly three to five minutes. Faster hardware would not change the verdicts, only the time taken.
- All submissions used are synthetic and short. A real tender pack is longer, more often scanned, and less tidily written.
- Only one checklist of ten items was used. A different checklist, or one written in different language, may produce different results.
- The optical character recognition step promised in the architecture document has not been built, so scanned documents cannot be checked at all yet. The system refuses them clearly rather than pretending to read them.

- These results describe one specific model at one specific setting. They are a starting point to improve on, not a final measure of what is achievable.