

Evaluation Results

Public Procurement Document-Completeness Agent

Model: llama3.1:8b | Prompt: v2.0-per-item | Method: one check per checklist item | Review threshold: 0.85 | Context: 16,384 tokens | 15 September 2026, 20:00 UTC

1 of 1 cases behaved as expected.

Tested against: A seven-page synthetic tender submission, the same content again as a Word file and as plain text, a scanned image-only version of it, a realistic tender package containing none of the required documents, and a sixty-page document used to test the size limit. Every document is invented; no real bidder or tender data was used.

By area

Area tested	As expected
Handling different file types	1 of 1

Every case

Case	Scenario	Expected	Actual	Result
EV-16	Same submission content supplied as a plain .txt file.	Found, with page_number 1 and a verbatim snippet	synthetic-submission.txt: Found (confidence 0.95, no page) "A bid guarantee in the sum of UGX 24,000,000.00 has been issued in fa...	Pass

Limitations

- The machine used for testing has no graphics card, so the model runs on the processor alone. Each check takes roughly three to five minutes. Faster hardware would not change the verdicts, only the time taken.
- All submissions used are synthetic and short. A real tender pack is longer, more often scanned, and less tidily written.
- Only one checklist of ten items was used. A different checklist, or one written in different language, may produce different results.
- The optical character recognition step promised in the architecture document has not been built, so scanned documents cannot be checked at all yet. The system refuses them clearly rather than pretending to read them.
- These results describe one specific model at one specific setting. They are a starting point to improve on, not a final measure of what is achievable.